



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Regularization

K. Breininger, V. Christlein, Z. Yang, L. Rist, M. Nau, S. Jaganathan, C. Liu, N. Maul, L. Folle, M. Zinnen,  
K. Packhäuser

Pattern Recognition Lab, Friedrich-Alexander University of Erlangen-Nürnberg

April 20, 2024



## Tasks in this exercise

1. Optimization Constraints: Augmenting the loss function
2. Dropout **Layer**
3. Batch Normalization **Layer**
4. LeNet: Put everything together (**optional**)
5. RNN layer: Elman Unit
6. LSTM layer: Backpropagation at its best! (**optional**)



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Optimization Constraints: Loss function augmentation



## General outline

- Constraints change the total loss ...
- ... and have influence on the weight update of the respective layer!

## General outline

- Constraints change the total loss ...
  - ... and have influence on the weight update of the respective layer!
  - Implement constraints as separate classes
- **Independent** of loss function

## General outline

- Constraints change the total loss ...
- ... and have influence on the weight update of the respective layer!
- Implement constraints as separate classes
- **Independent** of loss function
- Constraints **only need** current weights
- Add constraint objects in the optimizer

## General outline

- Constraints change the total loss ...
- ... and have influence on the weight update of the respective layer!
- Implement constraints as separate classes
- **Independent** of loss function
- Constraints **only need** current weights
- Add constraint objects in the optimizer
- Since constraints generate part of the loss:
- Change Neural Network container class (and associated classes) to “channel” and gather **regularization loss** for **all layers**

## Workflow

- Forward pass
- Calculate norm of weights in each trainable layer and gather as regularization loss
- Add regularization loss to the final loss



## Workflow

- Forward pass
  - Calculate norm of weights in each trainable layer and gather as regularization loss
  - Add regularization loss to the final loss
- Backward pass
  - In each trainable layer, include **the gradient of norm** when calculating update

## $L_2$ regularization

- Forward pass:

$$\tilde{L}(\mathbf{w}) = L(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$$

- Backward pass:

$$\mathbf{w}^{(k+1)} = \underbrace{(1 - \eta \lambda) \mathbf{w}^{(k)}}_{\text{Shrinkage}} - \eta \frac{\partial L}{\partial \mathbf{w}^{(k)}}$$

## $L_2$ regularization

- Forward pass:

$$\tilde{L}(\mathbf{w}) = L(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$$

- Backward pass:

$$\mathbf{w}^{(k+1)} = \underbrace{(1 - \eta \lambda)}_{\text{Shrinkage}} \mathbf{w}^{(k)} - \eta \frac{\partial L}{\partial \mathbf{w}^{(k)}}$$

- In the Forward pass the  $L_2$  norm gets squared, which eliminates the square root inside and increases the numerical stability as the gradient is easier to compute.

## $L_2$ regularization

- Forward pass:

$$\tilde{L}(\mathbf{w}) = L(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$$

- Backward pass:

$$\mathbf{w}^{(k+1)} = \underbrace{(1 - \eta \lambda)}_{\text{Shrinkage}} \mathbf{w}^{(k)} - \eta \frac{\partial L}{\partial \mathbf{w}^{(k)}}$$

- In the Forward pass the  $L_2$  norm gets squared, which eliminates the square root inside and increases the numerical stability as the gradient is easier to compute.
- Notice for matrices we compute here the Frobenius norm, not the Spectral norm.

## $L_2$ regularization

- Forward pass:

$$\tilde{L}(\mathbf{w}) = L(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$$

- Backward pass:

$$\mathbf{w}^{(k+1)} = \underbrace{(1 - \eta \lambda) \mathbf{w}^{(k)}}_{\text{Shrinkage}} - \eta \frac{\partial L}{\partial \mathbf{w}^{(k)}}$$

- In the Forward pass the  $L_2$  norm gets squared, which eliminates the square root inside and increases the numerical stability as the gradient is easier to compute.
- Notice for matrices we compute here the Frobenius norm, not the Spectral norm.
- The influence of constraints is controlled via  $\lambda$ . Because `lambda` is a python keyword, you want to use e.g. `alpha` instead.

## $L_1$ regularization

- Forward pass:

$$\tilde{L}(\mathbf{w}) = L(\mathbf{w}) + \lambda \|\mathbf{w}\|_1$$

- Backward pass:

$$\mathbf{w}^{(k+1)} = \underbrace{\mathbf{w}^{(k)} - \eta \lambda \text{sign}(\mathbf{w}^{(k)})}_{\text{Other shrinkage}} - \eta \frac{\partial L}{\partial \mathbf{w}^{(k)}}$$



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Dropout



## Method




Figure: Dropout

- Implement this as a **fixed-function layer**



## Method




Figure: Dropout

- Implement this as a **fixed-function layer**
- Randomly set **activations**  $\mapsto 0$  with probability  $1 - p$

## Method

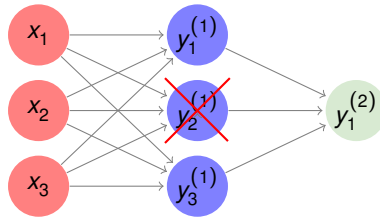


Figure: Dropout

- Implement this as a **fixed-function layer**
- Randomly set **activations**  $\mapsto 0$  with probability  $1 - p$
- **Test-time**: multiply activations with  $p$

## Inverted Dropout

- Can we get rid of the dropout layer at test-time?

## Inverted Dropout

- Can we get rid of the dropout layer at test-time?
- Change the behavior during training
- Multiply activations in forward-pass **only during training** by  $\frac{1}{p}$
- Note: the backward pass has to be adapted as well!



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

# Batch normalization



## Forward pass

→ Normalization as a new layer with 2 parameters,  $\gamma$  and  $\beta$

## Forward pass

→ Normalization as a new layer with 2 parameters,  $\gamma$  and  $\beta$

$$\tilde{\mathbf{X}} = \frac{\mathbf{X} - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

$\mu_B$  and  $\sigma_B^2$  from **batch**

## Forward pass

→ Normalization as a new layer with 2 parameters,  $\gamma$  and  $\beta$

$$\tilde{\mathbf{X}} = \frac{\mathbf{X} - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

$\mu_B$  and  $\sigma_B^2$  from **batch**

$$\hat{\mathbf{Y}} = \gamma \tilde{\mathbf{X}} + \beta$$



## Forward pass

→ Normalization as a new layer with 2 parameters,  $\gamma$  and  $\beta$

$$\tilde{\mathbf{X}} = \frac{\mathbf{X} - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

$\mu_B$  and  $\sigma_B^2$  from **batch**

$$\hat{\mathbf{Y}} = \gamma \tilde{\mathbf{X}} + \beta$$

- $\mu, \sigma^2$  have the **same dimension** as the **input vectors**

## Forward pass

→ Normalization as a new layer with 2 parameters,  $\gamma$  and  $\beta$

$$\tilde{\mathbf{X}} = \frac{\mathbf{X} - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

$\mu_B$  and  $\sigma_B^2$  from **batch**

$$\hat{\mathbf{Y}} = \gamma \tilde{\mathbf{X}} + \beta$$

- $\mu, \sigma^2$  have the **same dimension** as the **input vectors**
- $\beta, \gamma$  and  $\mu_B, \sigma_B^2$  have same **dimension** to be able to preserve **identity**

## Forward pass

→ Normalization as a new layer with 2 parameters,  $\gamma$  and  $\beta$

$$\tilde{\mathbf{X}} = \frac{\mathbf{X} - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

$\mu_B$  and  $\sigma_B^2$  from **batch**

$$\hat{\mathbf{Y}} = \gamma \tilde{\mathbf{X}} + \beta$$

- $\mu, \sigma^2$  have the **same dimension** as the **input vectors**
- $\beta, \gamma$  and  $\mu_B, \sigma_B^2$  have same **dimension** to be able to preserve **identity**
- Notice that  $\beta$  is a **bias**

## Test time

- Test-time: replace  $\mu_B$  and  $\sigma_B^2$  with  $\mu$  and  $\sigma^2$  of the **training set**

## Test time

- Test-time: replace  $\mu_B$  and  $\sigma_B^2$  with  $\mu$  and  $\sigma^2$  of the **training set**
- It's **expensive** to calculate the true training set mean and variance

## Test time

- Test-time: replace  $\mu_B$  and  $\sigma_B^2$  with  $\mu$  and  $\sigma^2$  of the **training set**
- It's **expensive** to calculate the true training set mean and variance
- Therefore a **moving average** is common:

$$\begin{aligned}\tilde{\mu}^{(k)} &\approx \alpha \tilde{\mu}^{(k-1)} + (1 - \alpha) \mu_B^{(k)} \\ \tilde{\sigma}^{2(k)} &\approx \alpha \tilde{\sigma}^{2(k-1)} + (1 - \alpha) \sigma_B^{2(k)}\end{aligned}$$

## Test time

- Test-time: replace  $\mu_B$  and  $\sigma_B^2$  with  $\mu$  and  $\sigma^2$  of the **training set**
- It's **expensive** to calculate the true training set mean and variance
- Therefore a **moving average** is common:

$$\begin{aligned}\tilde{\mu}^{(k)} &\approx \alpha \tilde{\mu}^{(k-1)} + (1 - \alpha) \mu_B^{(k)} \\ \tilde{\sigma}^{2(k)} &\approx \alpha \tilde{\sigma}^{2(k-1)} + (1 - \alpha) \sigma_B^{2(k)}\end{aligned}$$

- Moving average **decay**  $\alpha$  (e.g. 0.8)
- The exponent (k) and (k-1) are iteration-indices!

## Backward pass

- Gradient **with respect to weights** is simply:

$$\frac{\partial L}{\partial \gamma} = \sum_{b=1}^B \frac{\partial L}{\partial \hat{\mathbf{Y}}_b} \tilde{\mathbf{X}}_b = \sum_{b=1}^B \mathbf{E}_b \tilde{\mathbf{X}}_b$$

- For the **bias** likewise we have:

$$\frac{\partial L}{\partial \beta} = \sum_{b=1}^B \frac{\partial L}{\partial \hat{\mathbf{Y}}_b} = \sum_{b=1}^B \mathbf{E}_b$$



## Backward pass

The **gradient with respect to the input** is more complicated, but here it is:

$$\begin{aligned}
 \frac{\partial L}{\partial \tilde{\mathbf{X}}} &= \frac{\partial L}{\partial \hat{\mathbf{Y}}} \odot \gamma \\
 \frac{\partial L}{\partial \sigma_B^2} &= \sum_{b=1}^B \frac{\partial L}{\partial \tilde{\mathbf{X}}_b} \odot (\mathbf{x}_b - \mu_B) \odot \frac{-1}{2} (\sigma_B^2 + \epsilon)^{-\frac{3}{2}} \\
 \frac{\partial L}{\partial \mu_B} &= \left( \sum_{b=1}^B \frac{\partial L}{\partial \tilde{\mathbf{X}}_b} \odot \frac{-1}{\sqrt{\sigma_B^2 + \epsilon}} \right) + \underbrace{\frac{\partial L}{\partial \sigma_B^2} \odot \frac{\sum_{b=1}^B -2(\mathbf{x}_b - \mu_B)}{B}}_0 \\
 \frac{\partial L}{\partial \mathbf{X}} &= \frac{\partial L}{\partial \tilde{\mathbf{X}}} \odot \frac{1}{\sqrt{\sigma_B^2 + \epsilon}} + \frac{\partial L}{\partial \sigma_B^2} \odot \frac{2(\mathbf{x} - \mu_B)}{B} + \frac{\partial L}{\partial \mu_B} \odot \frac{1}{B}
 \end{aligned}$$

## Backward pass

- $\odot$  denotes an element-wise multiplication. Always check the dimensionality of your matrices!

## Backward pass

- $\odot$  denotes an element-wise multiplication. Always check the dimensionality of your matrices!
- To make life easier, we will provide the code for the computation of the gradient with respect to the input:

## Backward pass

- $\odot$  denotes an element-wise multiplication. Always check the dimensionality of your matrices!
- To make life easier, we will provide the code for the computation of the gradient with respect to the input:
- `compute_bn_gradients`

## Convolutional Batch Normalization

- In CNNs batch normalization is adjusted to work **similar to convolution**

## Convolutional Batch Normalization

- In CNNs batch normalization is adjusted to work **similar to convolution**
- A scalar  $\mu, \sigma$  is calculated for the  $H$  **channels**

## Convolutional Batch Normalization

- In CNNs batch normalization is adjusted to work **similar to convolution**
- A scalar  $\mu, \sigma$  is calculated for the  $H$  **channels**
- Implementation can be reused, by observing
  - that **spatial dimensions**  $M, N$  can be treated **like the batch dimension**  $B$

## Convolutional Batch Normalization

- In CNNs batch normalization is adjusted to work **similar to convolution**
- A scalar  $\mu, \sigma$  is calculated for the  $H$  **channels**
- Implementation can be reused, by observing
  - that **spatial dimensions**  $M, N$  can be treated **like the batch dimension**  $B$
  - we can **reshape** the  $B \times H \times M \times N$  tensor to  $B \times H \times M \cdot N$
  - because of our format we have to **transpose** from  $B \times H \times M \cdot N$  to  $B \times M \cdot N \times H$
  - and afterwards **reshape again** to have a  $B \cdot M \cdot N \times H$  tensor



## Convolutional Batch Normalization

- In CNNs batch normalization is adjusted to work **similar to convolution**
- A scalar  $\mu, \sigma$  is calculated for the  $H$  **channels**
- Implementation can be reused, by observing
  - that **spatial dimensions**  $M, N$  can be treated **like the batch dimension**  $B$
  - we can **reshape** the  $B \times H \times M \times N$  tensor to  $B \times H \times M \cdot N$
  - because of our format we have to **transpose** from  $B \times H \times M \cdot N$  to  $B \times M \cdot N \times H$
  - and afterwards **reshape again** to have a  $B \cdot M \cdot N \times H$  tensor
- Consequently we have to **reverse this** before returning the **output**

## Convolutional Batch Normalization

- In CNNs batch normalization is adjusted to work **similar to convolution**
- A scalar  $\mu, \sigma$  is calculated for the  $H$  **channels**
- Implementation can be reused, by observing
  - that **spatial dimensions**  $M, N$  can be treated **like the batch dimension**  $B$
  - we can **reshape** the  $B \times H \times M \times N$  tensor to  $B \times H \times M \cdot N$
  - because of our format we have to **transpose** from  $B \times H \times M \cdot N$  to  $B \times M \cdot N \times H$
  - and afterwards **reshape again** to have a  $B \cdot M \cdot N \times H$  tensor
- Consequently we have to **reverse this** before returning the **output**
- ... and do the **same** in the **backward pass**



FRIEDRICH-ALEXANDER-  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
SCHOOL OF ENGINEERING

## LeNet (optional)



# LeNet architecture

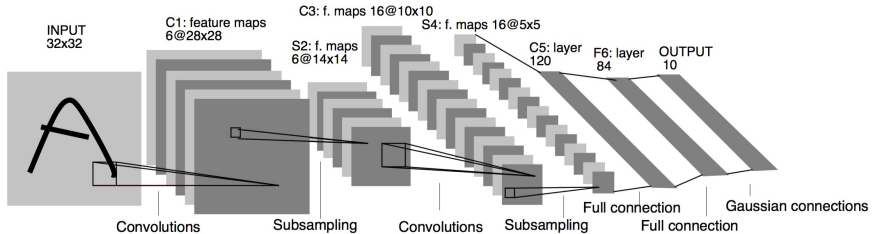


Figure: LeNet

# Modified LeNet architecture

## Deviations

- Input is  $28 \times 28$
- Our conv only supports “same” padding - so C3 has **larger activation maps**
- Input to **C5** is also **larger**
- We only implemented ReLUs, so **no** TanH
- We also use the implemented SoftMax **instead of** RBF units

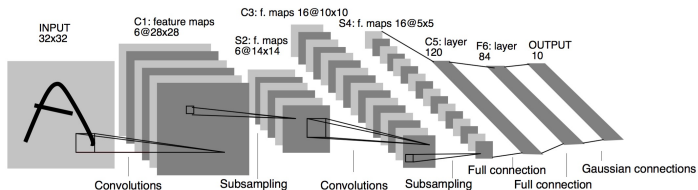


Figure: LeNet

Thanks for listening.  
**Any questions?**